

M4: ASH Parallel Design


Evan and Roy

ASH Overview + Sequential Methodology

- **Problem:** hyperparameters k^2 **model performance** – however, optimizing has become increasingly challenging b/c **more complex datasets** + **high-dimensional search spaces**
- **Solution:** Asynchronous Sequential Halving (ASH) is an algorithm for **efficiently tuning ML hyperparameters** by combing over the hyperparameter search space with **multiple cores simultaneously**
 - Random search > grid search b/c less exhaustive/expensive to get similar results (**~95%**)
- **Sequential Halving (SH):** sequentially iterate over all models on the base rung and promote the top half to the next rung, then double the epochs and iterate over all models in next rung and promote, etc.

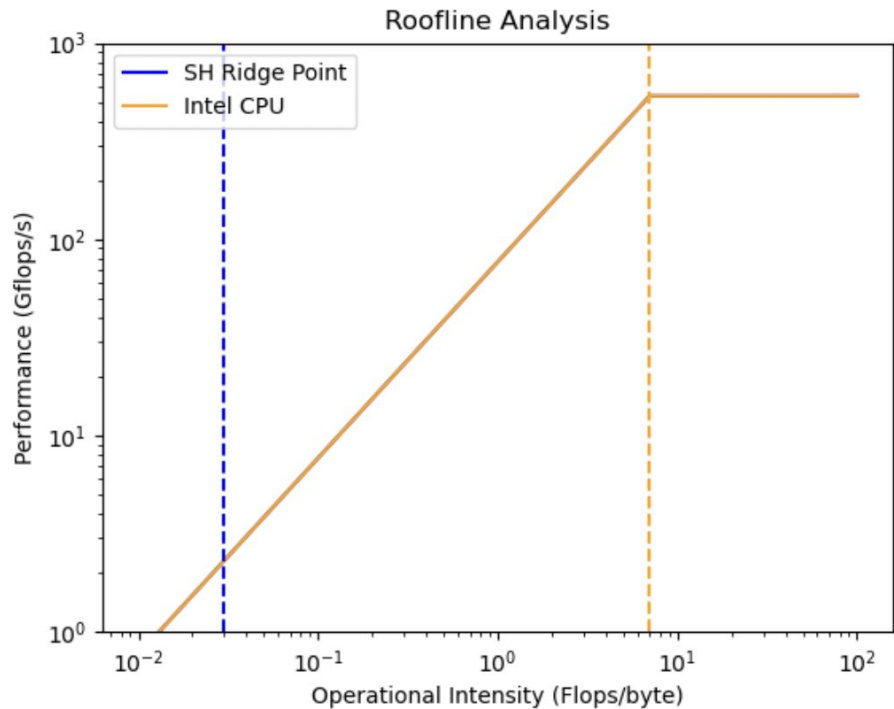
Sequential Baseline

Computation Time by Number of Models Broken Down by Rung



Roofline Analysis

- **Nominal Peak Performance** of Intel CPU: 537.6 Gflop/s
- **Beta**: 76.8 GB/s
- **Optimal Ridge Point**: 7 Flops/byte
- The operational intensity of our bottleneck: 7500 flops/250000 bytes = 0.03 flops/byte
- **Performance**: $76.8 * 0.03 = 2.304$ Gflop/s
- Gprof indicates **intensive memory management** and **movement**
- Supports hypothesis that increasing processes will be effective in **increasing operational intensity**

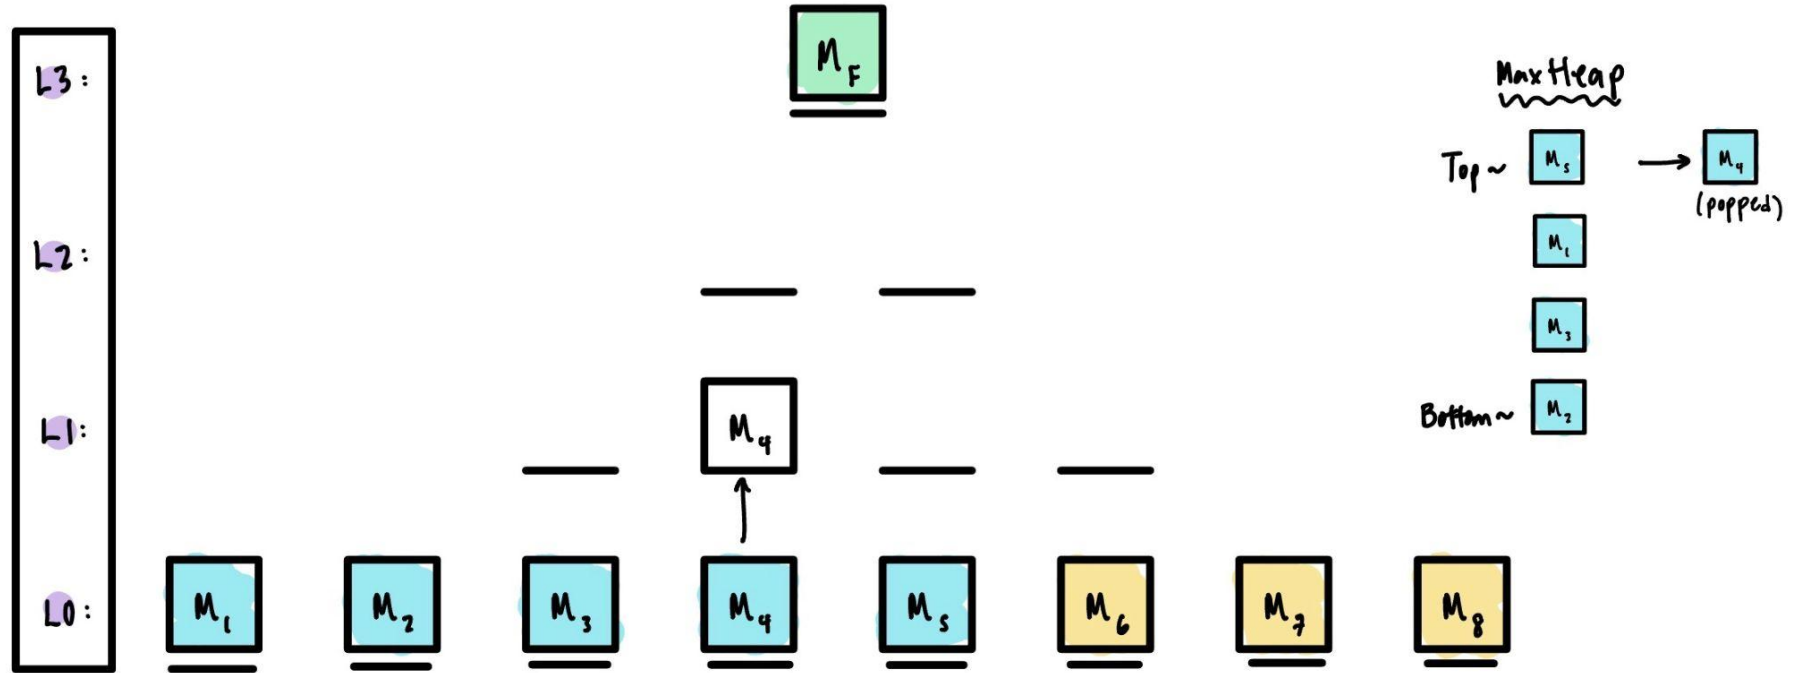


Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
8.98	0.35	0.35	145228405	0.00	0.00	std::vector<float, std::allocator<float> >::operator
8.72	0.69	0.34	95008211	0.00	0.00	std::vector<float, std::allocator<float> >::size() c
7.18	0.97	0.28	100210474	0.00	0.00	std::vector<std::vector<float, std::allocator<float
5.39	1.18	0.21	2	0.11	1.78	NeuralNetwork::train(std::vector<std::vector<float,

Parallel Design



Driver Code: MPI to 1) send model hyperparameter configurations to cores
2) communicate individual model results